

Docker, Git(hub)

MAS Departmental Disclaimers:

For students trying to take or audit from outside the MAS program.

Taking or auditing 400 courses is simply is not permitted because this is a self-supporting program. Sorry, unfortunately, you will NOT be able to take any of the 400 level Stats courses.

There are NO exceptions that can be made by the department. These classes were designed specifically for students who applied directly to the program.

The students of this program are also not allowed to audit or enroll in classes outside of the program as it was created for working professionals.

If you would like to apply for the program, you are welcome to do so: <https://master.stat.ucla.edu/admissions/>

Information is found here: <https://master.stat.ucla.edu/> And here: <https://master.stat.ucla.edu/faq/>



A reminder...

class communications

Slack Channel
uclastat418-class

Last time we saw an introduction to docker



today we'll make use of it while also hearing a bit more about what it is.

throughout the rest of course (in most weeks) we will continue to make use of docker environments to incrementally learn about its features

Docker: what is it?

a software to build and run containers with a container being a single instance of an image

an image is a unit of software that allows one to package up code and all its dependencies

Docker: why is it important?

different software development stacks require different environments;
especially dependent on task

classical solutions entail virtual environments (anaconda, python
virtualenv) or virtual machines (vagrant, vmware)

Docker: so why do data science with it?

it is platform agnostic; as long as a machine has docker we can use the same environment (perfect in a classroom setting where we all have different machines; perfect in a work setting where we might have different machines and certainly different versions)

can share version-controlled environments akin to GitHub for docker images

containers are ephemeral; if you mess up you can start over easily

build on top of other peoples work

ready-made workflow for deployment through Kubernetes or similar

More on Containerization

Containerization involves packaging software code along with the essential operating system libraries and dependencies it needs to run, creating a self-contained and lightweight executable called a container. This container is designed to run consistently across any infrastructure, eliminating the inconsistencies often encountered when moving applications between different computing environments.

Compared to traditional virtual machines, containers offer greater portability and are more resource-efficient, making them the standard building blocks for modern cloud-native applications. This technology empowers developers to build and deploy applications more rapidly and securely by bundling all necessary components, ensuring they run reliably regardless of the underlying platform or cloud environment – effectively enabling a "write once, run anywhere" approach. ready-made workflow for deployment through Kubernetes or similar

While the underlying concept has existed for decades, the widespread adoption of containerization was significantly accelerated by the emergence of Docker in 2013. Docker provided user-friendly tools and a standardized way to package applications, leading organizations to increasingly leverage containers for developing new applications and modernizing existing ones for the cloud, benefiting from their lightweight nature, faster startup times, and the ability to achieve higher server utilization and reduced costs.

How Containerization Works

Encapsulating Applications - Containers package an application as one executable. This bundle includes the application code, configuration files, libraries, and dependencies needed to run.

Sharing the Operating System - Unlike VMs, containers don't include a full OS. Instead, they rely on a container runtime (like Docker) installed on the host OS. This runtime allows containers to share the host's operating system.

Efficiency Through Sharing and Isolation - Containers can share common resources like binaries and libraries. This sharing reduces overhead and makes containers smaller and faster to start than VMs. The isolation of each container also enhances security, preventing issues in one container from affecting others or the host system.

Achieving Portability and Consistency - By abstracting away from the host OS, containerized applications become highly portable. They can run uniformly across various platforms and clouds, from desktops to VMs, and across different operating systems or server environments.

Enabling Faster and Secure Development - Containerization speeds up and secures application development and deployment. This applies to both traditional monolithic applications and modern microservices architectures. Developers can build new cloud applications as containerized microservices or repackaging existing applications for better resource use.

Virtualization vs Containerization

Both virtualization and containerization improve computing efficiency by enabling multiple software instances to run on a single physical machine.

Virtualization: Bundling the OS

- Utilizes a hypervisor to separate the OS and applications from the underlying hardware.
- Allows running multiple operating systems (e.g., Windows and Linux) and applications simultaneously on one host.
- Each VM includes its own operating system, application files, libraries, and dependencies.
- Achieves cost savings by consolidating multiple workloads onto fewer physical servers.

Containerization: Sharing the OS

- Packages applications with their necessary dependencies into a single executable.
- Does not include a separate operating system. Instead, containers share the host OS through a container runtime.
- Often described as "lightweight" due to OS sharing and smaller size.
- Offers faster startup times and allows more applications to run on the same hardware compared to VMs.
- Leads to even higher server efficiencies and reduced costs.

Docker: running Rstudio

run the following

```
docker run -e PASSWORD='class' --rm -p 8787:8787 rocker/rstudio
```

and then visit in any browser

<http://localhost:8787> or <http://<your-ip-address>:8787>

finally login with the following credentials

Username: rstudio Password: <what you used above>

Docker: running Rstudio

to exit



We will see another way to exit in a moment

Docker: running Rstudio with mounting local volume (connect to your computer)

cd to desired directory, then use for Mac

```
docker run -d --rm -e PASSWORD=class -p 8787:8787 -v `pwd`: /home/rstudio/Documents  
rocker/rstudio
```

chdir to desired directory in Windows command line

```
docker run -d --rm -e PASSWORD=class -p 8787:8787 -v %cd%: /home/rstudio/Documents  
rocker/rstudio
```

cd to desired directory in Windows powershell

```
docker run -d --rm -e PASSWORD=class -p 8787:8787 -v ${PWD}: /home/rstudio/Documents  
rocker/rstudio
```

Again visit in your browser port 8787 either at local host or <your-ip-address> and login with your credentials

Docker: a few other commands

Now to exit we will need to take a look at what containers are running to find the container id or name

```
docker container ls
```

```
docker kill <container_name>
```

And finally, to see what docker images we have available

```
docker images
```

Docker: a few other commands

And finally, to see what docker images exist on our machine

```
docker images
```

If you'd like to remove any of these images

```
docker rmi <image id>
```


Git

Git: what is it?

a version control system to track the history of changes as people and teams collaborate on code projects together

Git manages the evolution of a set of files (or in our case a data analysis or model development) in a repository


this provides a transparent history of changes, who made them and how each

 Git turns 20 today!

To celebrate, we're sharing our exclusive interview with creator Linus Torvalds on its origins and impact.

Git revolutionized developer collaboration, and we at G... [about.gitlab.com/blog/2.](https://about.gitlab.com/blog/2025/04/07/git-turns-20/)

50+ likes

 gitlab
X (Twitter) · 4 hours ago



04/07/2025

A history

In 2005, Linus Torvalds was developing Linux when BitKeeper (proprietary source-control management) revoked their free license. Torvalds set 4 goals for his development needs for a SCM and when he couldn't find anything sufficient he decided to create his own. Development began April 3 and was released April 7.

"I'm an egotistical bastard, and I name all my projects after myself. First 'Linux', now 'Git'."

('git' is British slang for "pig headed, think they are always correct, argumentative").
















[Public](#)
[Watch 2417](#)
[Fork 26.1k](#)
[Star 54.2k](#)

e83c516 6 Branches 974 Tags

Go to file

<> Code ▾

 Linux Torvalds	Initial revision of "git", the information manager from hell	e83c516 · 21 years ago	 1 Commit
 Makefile	Initial revision of "git", the information manager from hell	21 years ago	
 README	Initial revision of "git", the information manager from hell	21 years ago	
 cache.h	Initial revision of "git", the information manager from hell	21 years ago	
 cat-file.c	Initial revision of "git", the information manager from hell	21 years ago	
 commit-tree.c	Initial revision of "git", the information manager from hell	21 years ago	
 init-db.c	Initial revision of "git", the information manager from hell	21 years ago	
 read-cache.c	Initial revision of "git", the information manager from hell	21 years ago	
 read-tree.c	Initial revision of "git", the information manager from hell	21 years ago	
 show-diff.c	Initial revision of "git", the information manager from hell	21 years ago	
 update-cache.c	Initial revision of "git", the information manager from hell	21 years ago	
 write-tree.c	Initial revision of "git", the information manager from hell	21 years ago	

 README

GIT – the stupid content tracker

"git" can mean anything, depending on your mood.







- random three-letter combination that is pronounceable, and not actually used by any common UNIX command. The fact that it is a mispronunciation of "get" may or may not be relevant.
- stupid. contemptible and despicable. simple. Take your pick from the dictionary of slang.
- "global information tracker": you're in a good mood, and it actually works for you. Angels sing, and a light suddenly fills the room.
- "goddamn idiotic truckload of sh*t": when it breaks

This is a stupid (but extremely fast) directory content manager. It doesn't do a whole lot, but what it does do is track directory contents efficiently.

About

Git Source Code Mirror - This is a publish-only repository but pull requests can be turned into patches to the mailing list via GitGitGadget (<https://gitgitgadget.github.io/>). Please follow Documentation/SubmittingPatches procedure for any of your improvements.

- c shell hacktoberfest

-  Readme
-  View license
-  Code of conduct
-  Security policy
-  Activity
-  Custom properties

☆ 54.2k stars

👁 2.4k watching

 26.1k forks

Report repository


Releases

 974 tags

Packages

No packages published

Contributors 1,746



+ 1,732 contributors

<https://github.com/git/git/tree/e83c5163316f89bfbde7d9ab23ca2e25604af290>

Git

I< < PREV RANDOM NEXT > >I

THIS IS GIT. IT TRACKS COLLABORATIVE WORK
ON PROJECTS THROUGH A BEAUTIFUL
DISTRIBUTED GRAPH THEORY TREE MODEL.

COOL. HOW DO WE USE IT?

NO IDEA. JUST MEMORIZE THESE SHELL
COMMANDS AND TYPE THEM TO SYNC UP.
IF YOU GET ERRORS, SAVE YOUR WORK
ELSEWHERE, DELETE THE PROJECT,
AND DOWNLOAD A FRESH COPY.



I< < PREV RANDOM NEXT > >I

PERMANENT LINK TO THIS COMIC: [HTTPS://XKCD.COM/1597/](https://xkcd.com/1597/)

IMAGE URL (FOR HOTLINKING/EMBEDDING): [HTTPS://IMGS.XKCD.COM/COMICS/GIT.PNG](https://imgs.xkcd.com/comics/git.png)

Repository management services



Repository management services



GitHub ✓

Code Collaboration & Version Control

Follow

Stacks
32.5K

I Use This

Fans	Jobs	Votes
28.7K	3.75K	10K



Bitbucket ✓

Code Collaboration & Version Control

Follow

Stacks
9.82K

I Use This

Fans	Jobs	Votes
8.44K	438	2.77K



GitLab ✓

Code Collaboration & Version Control




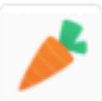


























Follow

Stacks
8.35K

I Use This

Fans	Jobs	Votes
7.57K	357	2K

Repository management services

What companies use GitHub?	What companies use Bitbucket?	What companies use GitLab?
5837 companies on StackShare use GitHub	2270 companies on StackShare use Bitbucket	1584 companies on StackShare use GitLab
<div><div></div><div>Airbnb</div></div> <div><div></div><div>Netflix</div></div> <div><div></div><div>Medium</div></div> <div><div></div><div>Instacart</div></div> <div><div></div><div>reddit</div></div> <div><div></div><div>Lyft</div></div> <div><div></div><div>StackShare</div></div> <div><div></div><div>Shopify</div></div> <div><div></div><div>9GAG</div></div> <div><div></div><div>Asana</div></div>	<div><div></div><div>PayPal</div></div> <div><div></div><div>Salesforce</div></div> <div><div></div><div>Zillow</div></div> <div><div></div><div>Starbucks</div></div> <div><div></div><div>CircleCI</div></div> <div><div></div><div>Tesla Motors</div></div> <div><div></div><div>Bitbucket</div></div> <div><div></div><div>Pandora</div></div> <div><div></div><div>Sellsuki</div></div> <div><div></div><div>Movielala</div></div>	<div><div></div><div>Alibaba.com</div></div> <div><div></div><div>Coderus</div></div> <div><div></div><div>Webedia</div></div> <div><div></div><div>Ticketmaster</div></div> <div><div></div><div>WebbyLab</div></div> <div><div></div><div>Edify</div></div> <div><div></div><div>Electronic Arts</div></div> <div><div></div><div>Freelancer.com</div></div> <div><div></div><div>Citrix</div></div> <div><div></div><div>WILD</div></div>

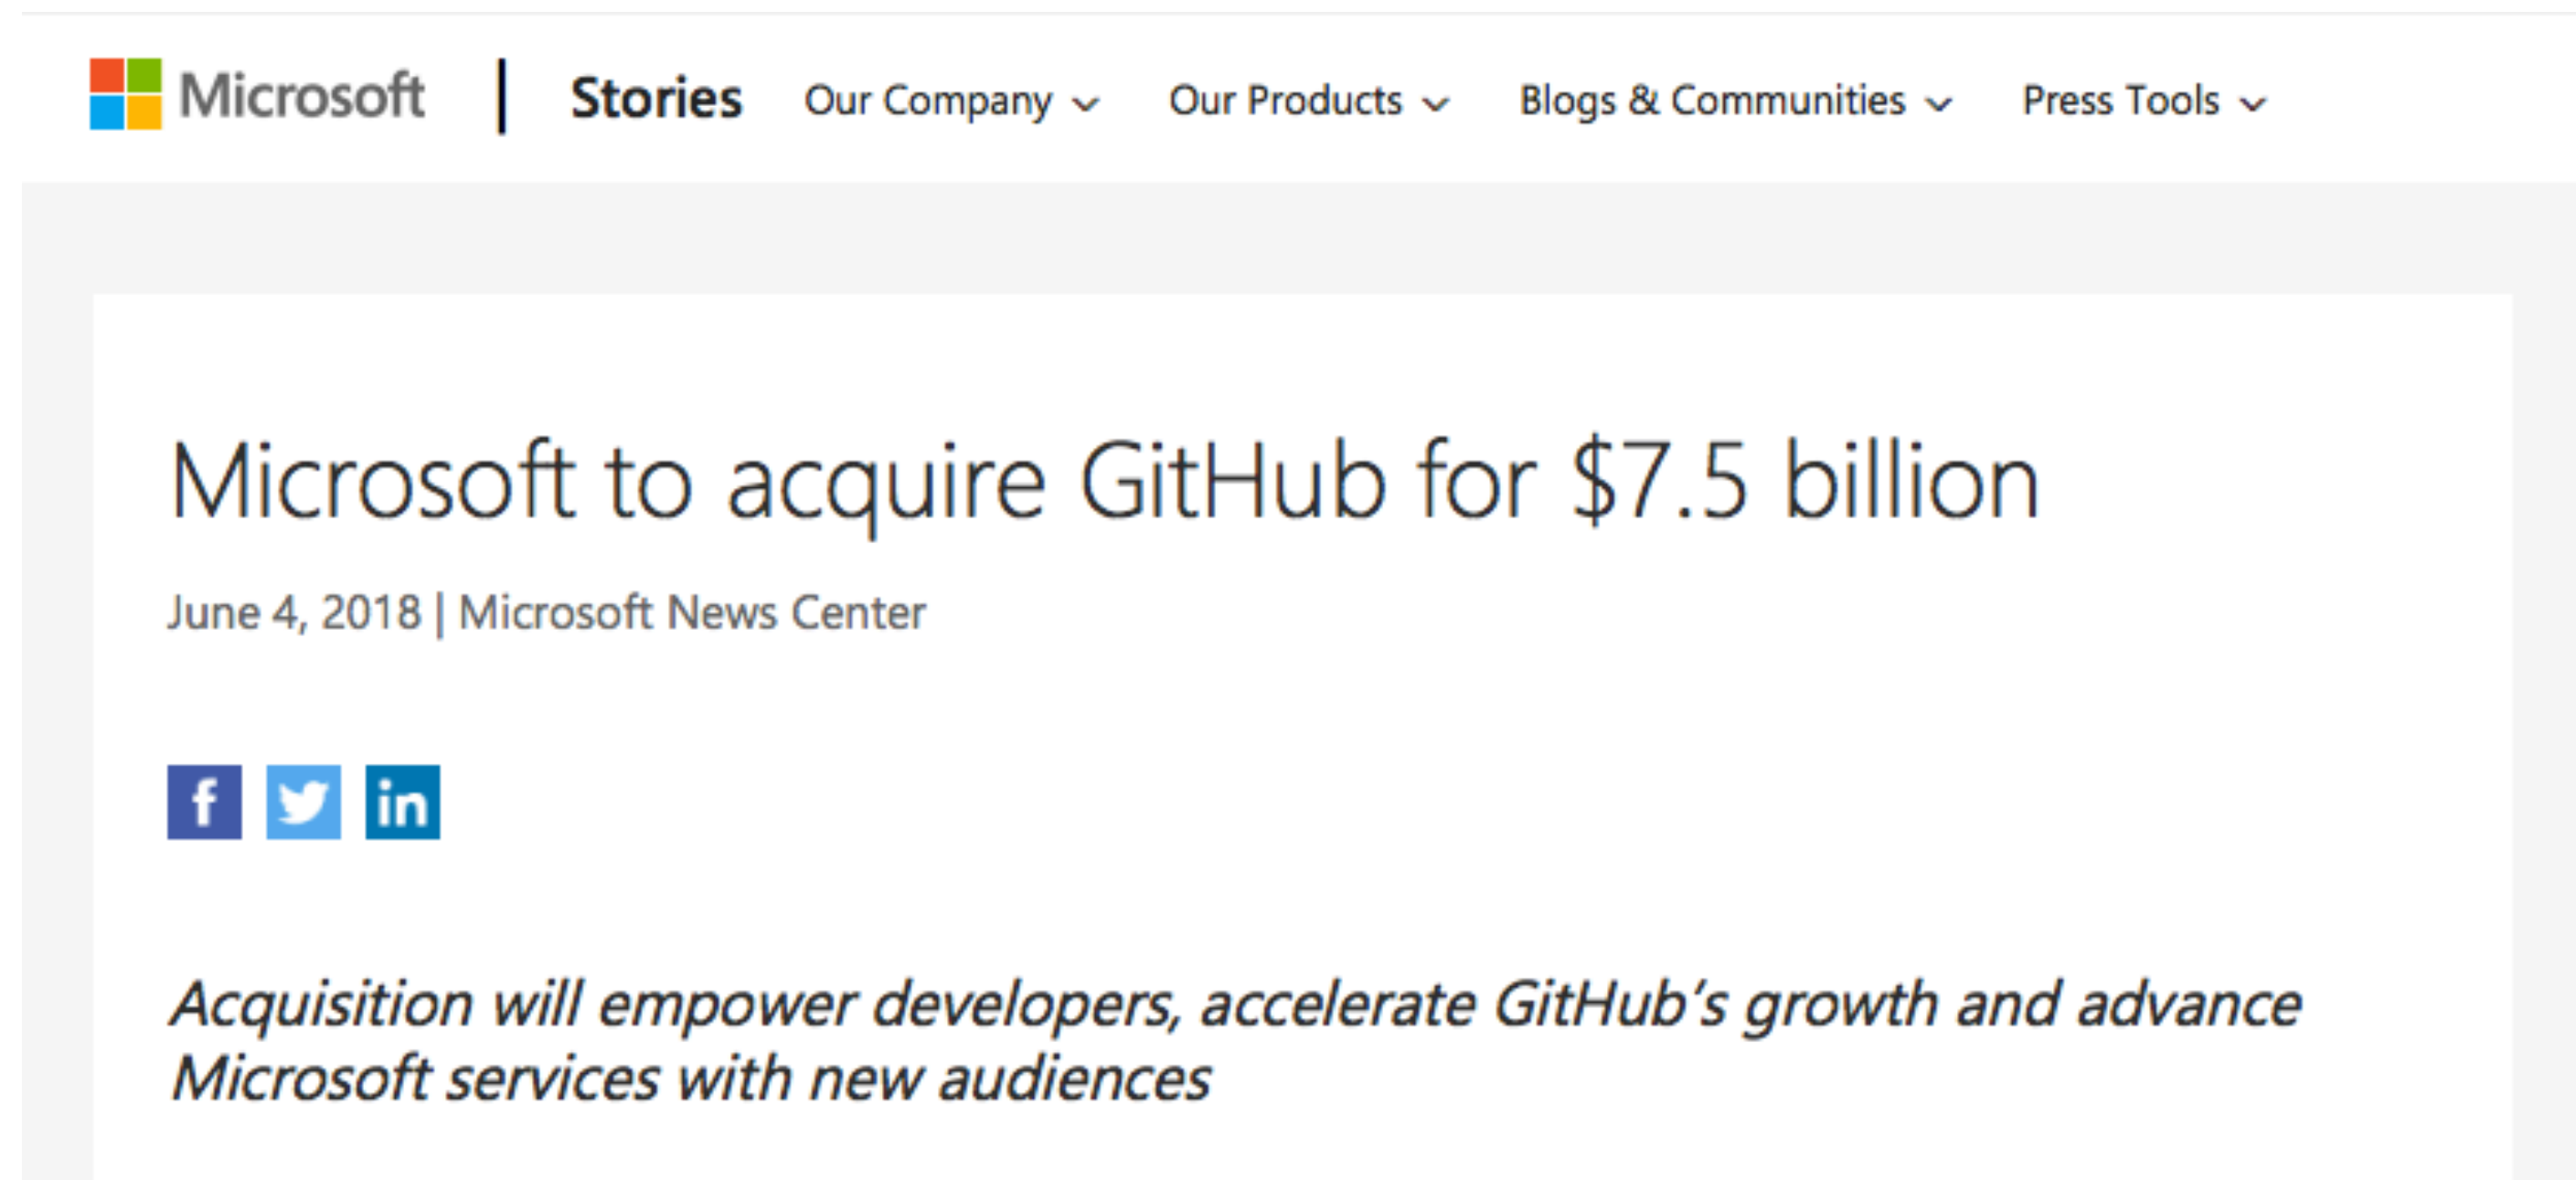


We're going to use GitHub


<http://github.com>

Make an account if you don't have one

Let it be known...



<https://www.theverge.com/2018/6/18/17474284/microsoft-github-acquisition-developer-reaction>



Mike Coutermarsh

@mscccc



Just googled "undo last commit but keep changes"

And I work for GitHub.

This stuff is hard!

11:10 PM - 18 Jan 2018

321 Retweets 1,007 Likes



39 321 1.0K

It's not all fun and games. Actually it's quite hard. Initially just working on solo repositories becomes fine but once collaborating you'll start to run into merge conflicts (yikes)

Keep going!!

GitHub: what tool to use?

Use what makes you feel comfortable and efficient.

“

“I certainly know Git very well, and honestly think I’m far faster and more efficient in a Git GUI than I could possibly be in the command line – and I’m certainly not slow in the CLI.”

”

—Dan Clarke, blogger and co-organizer of .Net Oxford group

“

“I sometimes encounter people who feel it’s “better” to use command line Git, but for very ill-defined reasons. These people may feel like they should work in the shell, even if it leads to Git-avoidance, frequent mistakes, or limiting themselves to a small set of ~3 Git commands. This is counterproductive.”

”

—Jenny Bryan, Software Engineer at RStudio & Adjunct Professor at the University of British Columbia

Both quotes take from <https://blog.axosoft.com/git-gui-vs-cli/>

GitHub GUIs



There are a wide variety of choices for GUIs if you are interested.

GitHub: some (basic) guidelines

Personally, I start all my repositories remotely (on [GitHub.com](https://github.com)). I find it much easier to then pull to my local machine and to then make my first commit

This includes forking others' repos and then cloning to local

Create new branches to make edits (add new features) to the master code base

Commit early and often; typically don't push data (especially when large) and definitely not credentials

(these are all quite simple; you'll surely run into many issues and there are many more resources online that will help you, but don't be afraid to try some new things in this course to learn how it works)

GitHub: our class repository

Let's create a repo for you to use and push hw to my master repo. cd to the directory where you want to

Fork the class repo on GitHub. Clone your https version of the repo.

Now let's go to docker and do this in a (hopefully) more familiar environment...get an Rstudio container running with a local volume mounted. (if its not still running at the moment)

...will walk through as a class.

A bit more on reproducibility...

there is a bit of crisis taking place in regards to reproducibility as a whole in science fields. and this true of the data science field as well.

We've seen some of the tools to enable us to do this reproducible work, but what concepts should we be thinking of as we use them?

Repeatability

or test–retest reliability is the variation in measurements taken by a single person or instrument on the same item, under the same conditions, and in a short period of time. A less-than- perfect test–retest reliability causes test–retest variability.


Your systems must be repeatable and reliable, the same query on the same data should return the same results. Otherwise, none of this applies.



Low repeatability, Low accuracy



High repeatability, Low accuracy



High Repeatability, High accuracy

Reproducibility

an analysis is *reproducible* if there is a specific set of computational functions/analyses (usually specified in terms of code) that exactly reproduce all of the numbers in a published paper from raw data. It is now recognized that a critical component of the scientific process is that data analyses can be reproduced.

Replicability

but just because a study is reproducible does not mean that it is *replicable*. Replicability is stronger than *reproducibility*. A study is only *replicable* if you perform the exact same experiment (at least) twice, collect data in the same way both times, perform the same data analysis, and arrive at the same conclusions.

Stodden's Taxonomy of Reproducibility

Computational

when detailed information is provided about code, software, hardware and implementation details.




Empirical

when detailed information is provided about code, software, hardware and implementation details.



Statistical

when detailed information is provided about code, software, hardware and implementation details.



Computational reproducibility should be...

Reviewable

The descriptions of the research methods can be independently assessed and the results judged credible

Auditable

sufficient records (including data and software) have been archived so that the research can be defended later if necessary or differences between independent confirmations resolved. The archive might be private, as with traditional laboratory notebooks.

Replicable

tools are made available that would allow one to duplicate the results of the research, for example by running the authors' code to produce the plots shown in the publication.

Confirmable

the main conclusions of the research, or the model in production can be attained independently without the use of software provided by the author.

Simple rules for reproducible computational research

track results - whenever a result may be of potential interest, keep track of how it was produced. as a minimum, you should at least record sufficient details on programs, parameters, and manual procedures to allow yourself, in a year or so, to approximately reproduce the results.

script everything - whenever possible, rely on the execution of programs instead of manual procedures to modify data. If manual operations cannot be avoided, you should as a minimum note down which data files were modified or moved, and for what purpose.

store program versions - In order to exactly reproduce a given result, it may be necessary to use programs in the exact versions used originally. As a minimum, you should note the exact names and versions of the main programs you use.

Simple rules for reproducible computational research

use version control - even the slightest change to a computer program can have large intended or unintended consequences. As a minimum, you should archive copies of your scripts from time to time, so that you keep a rough record of the various states the code has taken during development

store data & intermediate results - in principle, as long as the full process used to produce a given result is tracked, all intermediate data can also be regenerated. In practice, having easily accessible intermediate results may be of great value. As a minimum, archive any intermediate result files that are produced when running an analysis

set a random number seed- many analyses and predictions include some element of randomness, meaning the same program will typically give slightly different results every time it is executed. As a minimum, note which analysis steps involve randomness, so that a level of discrepancy can be anticipated when reproducing the results.

Simple rules for reproducible computational research

store data viz inputs - from the time a figure is first generated to it being part of an analysis. It is critical to store the data and process that generated it. As a minimum, one should note which data formed the basis of a given plot and how this data could be reconstructed.

allow levels of analysis - in order to validate and fully understand the main result, it is often useful to inspect the detailed values underlying the summaries. Make those fluid and explorable, as a minimum at least once generate, inspect, and validate the detailed values underlying the summaries.

generate text programmatically- there is nothing quite as embarrassing as having a disagreement between your analytical narrative in a writeup and having the tables and figures disagree. Connect them so that there is no change of disagreement.

.